Represent Yourself - Product Specification

Executive Summary

Represent Yourself is an AI-powered legal assistant designed specifically for self-represented litigants (pro se / litigants in person). It combines document generation, case analysis, and procedural guidance into a native desktop application that helps ordinary people navigate complex legal proceedings without expensive legal representation.

Problem Statement

The Access to Justice Gap

1. **Legal representation is unaffordable** - Average hourly rates for lawyers range from $200-$500+. A typical civil case costs $10,000-$50,000+ in legal fees.

2. **Self-representation is overwhelming** - Court procedures, legal terminology, document formatting, and deadlines create massive barriers for non-lawyers.

3. **Existing tools are inadequate** - Legal templates are generic. ChatGPT gives generic answers. Neither understands your specific case context.

4. **Information asymmetry** - Opposing parties with lawyers have massive advantages in knowing procedures, deadlines, and effective arguments.

Solution

Represent Yourself provides:

- **Deep case understanding** - Upload all your documents once, AI builds comprehensive case profile
- **Document generation** - Court-ready documents with proper formatting and legal language
- **Procedural guidance** - Step-by-step guidance through court procedures
- **Strategic analysis** - Honest assessment of case strengths, weaknesses, and risks
- **Deadline management** - Never miss a filing deadline

- **Hearing preparation** - Mock hearings and cross-examination practice

Target Users

Primary: Self-Represented Litigants
- **Civil litigation** defendants/plaintiffs
- **Family law** (divorce, custody, child support)
- **Employment disputes** (wrongful termination, discrimination)
- **Landlord-tenant** disputes
- **Small claims** escalated to higher courts
- **Consumer protection** cases

User Personas

1. Sarah - Divorce Litigant - Going through contested divorce - Spouse has lawyer, she doesn't - Needs help with financial disclosures, custody arrangements - Budget: Can't afford $300/hour lawyer

2. Marcus - Employment Dispute - Wrongfully terminated, filed complaint - Company responded with motion to dismiss - Needs to file opposition within 14 days - Has evidence but doesn't know how to present it

3. Chen - Landlord Dispute - Suing landlord for security deposit - Case escalated beyond small claims - Intimidated by formal court procedures - English is second language
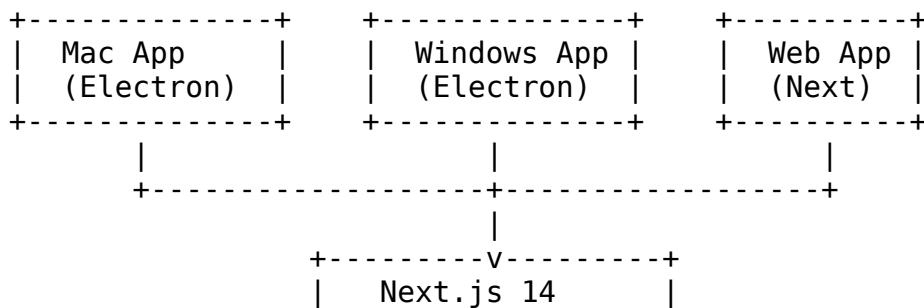
Secondary: Legal Aid Organizations
- Provide tool to clients they can't fully represent
- Training resource for pro bono attorneys
- Intake and case assessment assistance
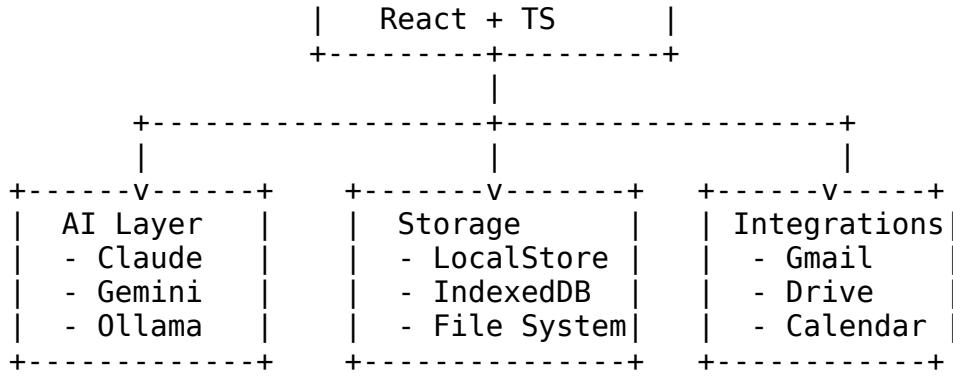
Product Architecture

Platform Strategy

```
                REPRESENT YOURSELF


    +--------------+   +--------------+   +----------+
    |  Mac App     |   |  Windows App |   | Web App  |
    |  (Electron)  |   |  (Electron)  |   | (Next)   |
    +--------------+   +--------------+   +----------+
           |                  |                |
        +------------------+-----------------+
                           |
        +---------v---------+
        |    Next.js 14     |
```

```
              |   React + TS      |
              +---------+---------+
                        |
        +---------------+-------------------+
        |               |                   |
  +------v------+  +-------v-------+  +------v-----+
  | AI Layer    |  |  Storage      |  | Integrations|
  |  - Claude   |  |  - LocalStore |  |  - Gmail    |
  |  - Gemini   |  |  - IndexedDB  |  |  - Drive    |
  |  - Ollama   |  |  - File System|  |  - Calendar |
  +-------------+  +---------------+  +------------+
```

Technical Stack

Layer	Technology	Purpose
Frontend	Next.js 14, React 18, TypeScript	App framework
Styling	Tailwind CSS, Radix UI	Design system
Desktop	Electron 28	Native Mac/Windows apps
State	Zustand + localStorage	Client-side persistence
AI	Claude API, Gemini API, Ollama	LLM providers
Integrations	Google APIs (Gmail, Drive, Calendar)	Productivity
PDF	pdf.js, pdf-lib	Document handling

Data Architecture

```
User Data (100% Client-Side)
|
+-- Case Profile
|    +-- Parties (plaintiffs, defendants, lawyers)
|    +-- Court details (case number, judge, jurisdiction)
|    +-- Procedural history
|    +-- Related proceedings
|
+-- Documents
|    +-- Uploaded files (PDF, DOCX, images)
|    +-- AI analysis results
|    +-- Generated drafts
|    +-- Filed documents
|
+-- Exhibits
|    +-- Exhibit metadata
|    +-- Bundles
|    +-- Page references
|
+-- Timeline
```

```
|   +-- Events (filings, hearings, deadlines)
|   +-- AI-extracted dates
|   +-- User-added events
|
+-- Configuration
    +-- AI provider settings
    +-- Google OAuth tokens
    +-- UI preferences
```

Feature Set

Phase 1: Core (Current)

1. Case Setup & Onboarding
- AI-powered setup wizard
- Upload documents, AI extracts case details
- Guided questions to fill gaps
- Auto-generates case profile

2. Document Drafting
- **5 document types:**
  - Affidavit/Affirmation
  - Written Submissions
  - Skeleton Arguments
  - Letter to Court
  - Response/Reply
- Step-by-step wizard interface
- Court-specific formatting
- Export to DOCX/PDF

3. Document Analysis
- Upload any legal document
- AI extracts:
  - Document type and purpose
  - Key facts and dates
  - Legal issues raised
  - Required responses
  - Deadlines triggered

4. Evidence Management
- Exhibit cataloging with metadata
- Bundle creation for hearings
- Page numbering and indexing

- Status tracking (draft to filed to admitted)

5. Timeline Builder
- Chronological case visualization
- Event categorization (court/opposing/yours/victory)
- Deadline highlighting
- AI-assisted date extraction from documents

6. AI Assistant Chat
- Context-aware legal Q&A
- Knows your case details
- Explains procedures in plain English
- Suggests next steps

7. Plain Language Translator
- Paste legal jargon, get plain English
- Key term glossary
- "What this means for you" summary
- Action items extraction

Phase 2: Strategic Tools

8. Case Analyzer
- Comprehensive strength/weakness assessment
- Risk scoring (0-100)
- Win probability estimation
- Strategic recommendations

9. Procedure Guide
- Court-specific procedure handbook
- Step-by-step filing instructions
- Form requirements
- Fee schedules

10. Discovery Helper
- Analyze discovery requests
- Identify responsive documents
- Draft objections
- Privilege log assistance

11. Motion Response Wizard
- Parse opposing motions
- Generate point-by-point responses
- Legal authority suggestions

- Filing checklist

12. Deadline Calculator
- Auto-calculate procedural deadlines
- Calendar integration
- Reminder notifications
- "Days remaining" tracking

Phase 3: Advanced Features

13. Hearing Simulator
- Mock hearing practice
- AI plays opposing counsel
- Cross-examination prep
- Judge Q&A simulation

14. Opponent Intelligence
- Track opposing party patterns
- Analyze their filing history
- Predict likely arguments
- Identify weaknesses

15. Settlement Calculator
- Model settlement scenarios
- Risk-adjusted valuations
- Negotiation ranges
- Walk-away points

16. Risk Scorecard
- Multi-factor risk assessment
- Litigation cost projections
- Outcome probability matrix
- Decision support

17. Legal GPS
- Visual case navigation
- "You are here" indicator
- Next steps guidance
- Milestone tracking

18. Bundle Generator
- Auto-generate hearing bundles
- Court-compliant formatting
- Index generation

- Paginated PDF output

Phase 4: Integrations

19. Gmail Integration
- Monitor case-related emails
- Auto-categorize by sender
- Extract deadlines from emails
- Draft responses

20. Google Drive Integration
- Sync documents to Drive
- Collaborative editing
- Version history
- Backup protection

21. Calendar Integration
- Sync deadlines and hearings
- Reminder scheduling
- Preparation timelines
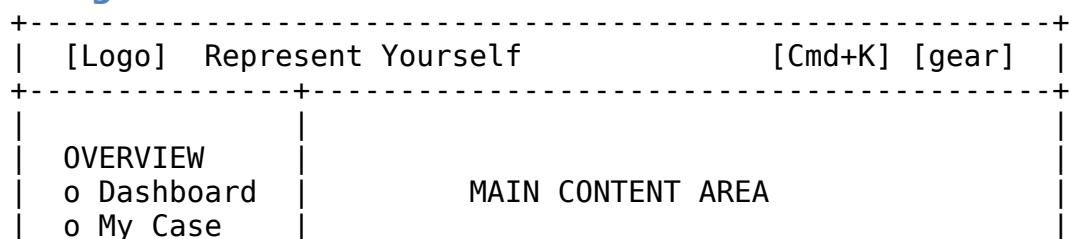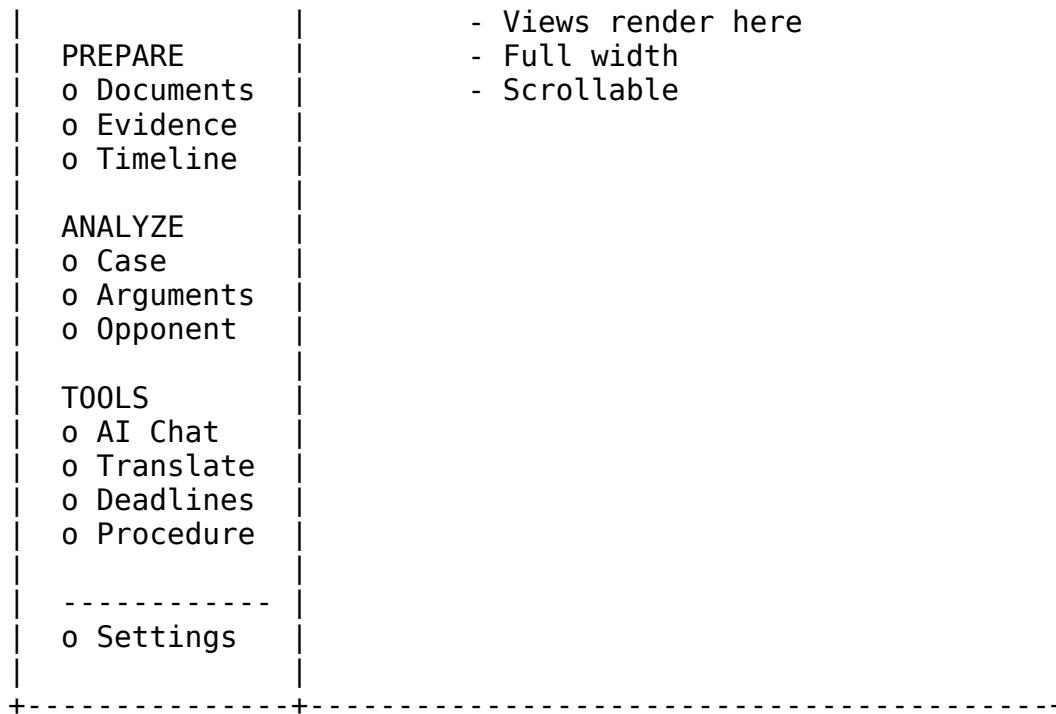
User Experience

Design Principles

1. **Calm, not overwhelming** - Legal matters are stressful. UI should feel supportive, not add anxiety.

2. **Guide, don't dump** - Never show 50 options at once. Progressive disclosure with clear next steps.

3. **Plain language first** - Explain everything in plain English. Legal terms always have definitions.

4. **Mobile-ready but desktop-focused** - Legal work needs big screens. Mobile for reference only.

5. **Offline-capable** - Documents and case data work offline. AI features need connection.

Navigation Structure

```
+--------------------------------------------------------+
|  [Logo]  Represent Yourself          [Cmd+K] [gear]  |
+---------------+----------------------------------------+
|               |                                        |
|  OVERVIEW     |                                        |
|  o Dashboard  |          MAIN CONTENT AREA             |
|  o My Case    |                                        |
```

```
|               |              - Views render here        |
|   PREPARE     |              - Full width                |
|   o Documents |              - Scrollable                |
|   o Evidence  |                                          |
|   o Timeline  |                                          |
|               |                                          |
|   ANALYZE     |                                          |
|   o Case      |                                          |
|   o Arguments |                                          |
|   o Opponent  |                                          |
|               |                                          |
|   TOOLS       |                                          |
|   o AI Chat   |                                          |
|   o Translate |                                          |
|   o Deadlines |                                          |
|   o Procedure |                                          |
|               |                                          |
|   ----------- |                                          |
|   o Settings  |                                          |
|               |                                          |
+---------------+------------------------------------------+
```
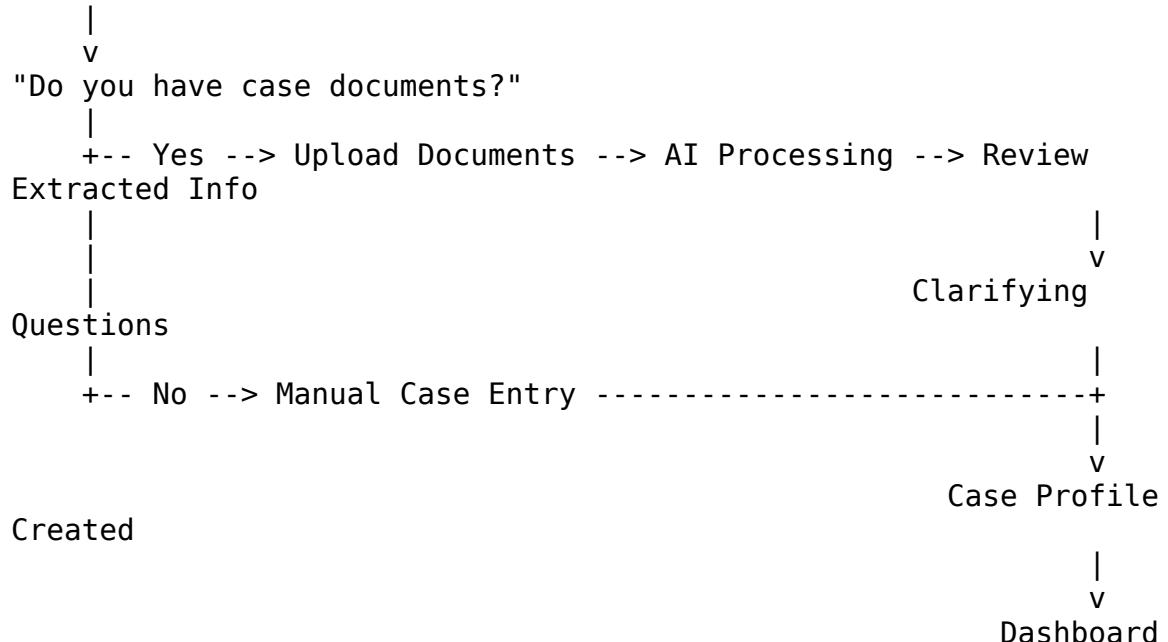
Key Flows
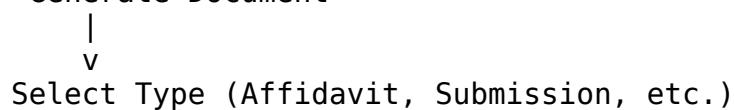
Flow 1: First-Time Setup

```
Welcome Screen
     |
     v
"Do you have case documents?"
     |
     +-- Yes --> Upload Documents --> AI Processing --> Review
Extracted Info
     |                                                 |
     |                                                 v
     |                                            Clarifying
Questions
     |                                                 |
     +-- No --> Manual Case Entry --------------------------+
                                                            |
                                                            v
                                                       Case Profile
Created
                                                            |
                                                            v
                                                       Dashboard
```
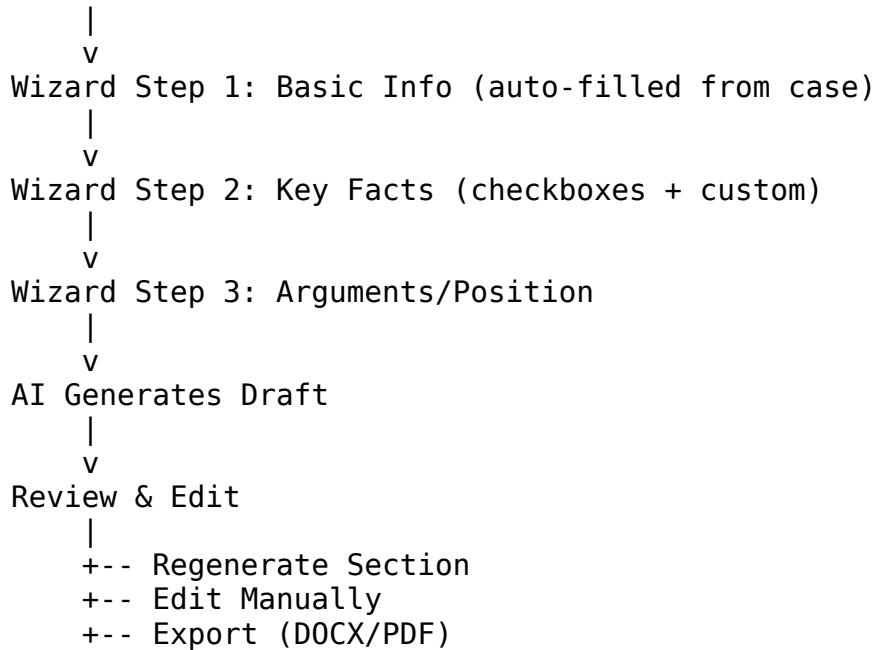
Flow 2: Document Generation

```
"Generate Document"
     |
     v
Select Type (Affidavit, Submission, etc.)
```

```
      |
      v
Wizard Step 1: Basic Info (auto-filled from case)
      |
      v
Wizard Step 2: Key Facts (checkboxes + custom)
      |
      v
Wizard Step 3: Arguments/Position
      |
      v
AI Generates Draft
      |
      v
Review & Edit
      |
      +-- Regenerate Section
      +-- Edit Manually
      +-- Export (DOCX/PDF)
```


AI Strategy

Strategic Foundation

Core Decision: API-First, Not Open Source

Open-source LLMs are a margin-expansion tool, not a foundation: - At low user counts, APIs are significantly cheaper - Fixed infrastructure costs for self-hosted: $3,500-8,000+/month - API costs at early stage: $400-900/month - Open-source only becomes economical at large, predictable scale

Why Not Self-Hosted Open Source: - High fixed infrastructure costs - DevOps + ML maintenance burden - Lower reasoning reliability for legal tasks - Silent failures are unacceptable in legal products - No frontier open model can run locally (Qwen3-235B needs 4-8x H100 GPUs)

Three-Layer Model Architecture

Users choose **quality tier**, not models. The system routes intelligently.

```
                  USER QUALITY SELECTION

   [ Standard ]        [ Professional ]        [ Premium ]
      Free                 $0.10/doc              $0.25/doc

                          |
                          ▼
```

```
┌────────────────────────────────────────────────────────┐
│            LAYER 1: INGESTION & STRUCTURING              │
│                    (Cheap Model)                         │
│                                                          │
│  • OCR and text extraction                               │
│  • Timeline extraction                                   │
│  • Entity recognition (parties, dates, amounts)          │
│  • Document classification                               │
│  • Metadata extraction                                   │
│                                                          │
│  Model: Gemini Flash / GPT-4o-mini                       │
│  Cost: ~$0.001-0.005 per document                        │
└────────────────────────────────────────────────────────┘
                            │
                            ▼
┌────────────────────────────────────────────────────────┐
│           LAYER 2: LEGAL REASONING & DRAFTING            │
│                    (Premium Model)                       │
│                                                          │
│  • Document generation (affidavits, submissions)         │
│  • Case strength analysis                                │
│  • Argument identification                               │
│  • Strategic recommendations                             │
│  • Motion response drafting                              │
│                                                          │
│  Standard: Claude Sonnet (default)                       │
│  Premium: Claude Opus (high-stakes toggle)               │
│  Cost: $0.01-0.10 per generation                         │
└────────────────────────────────────────────────────────┘
                            │
                            ▼
┌────────────────────────────────────────────────────────┐
│           LAYER 3: VERIFICATION & CONSISTENCY            │
│                    (Cheap Model)                         │
│                                                          │
│  • Citation checking                                     │
│  • Internal consistency validation                       │
│  • Format compliance                                     │
│  • Deadline verification                                 │
│  • Cross-reference validation                            │
│                                                          │
│  Model: Gemini Flash / GPT-4o-mini                       │
│  Cost: ~$0.001-0.005 per check                           │
└────────────────────────────────────────────────────────┘
```
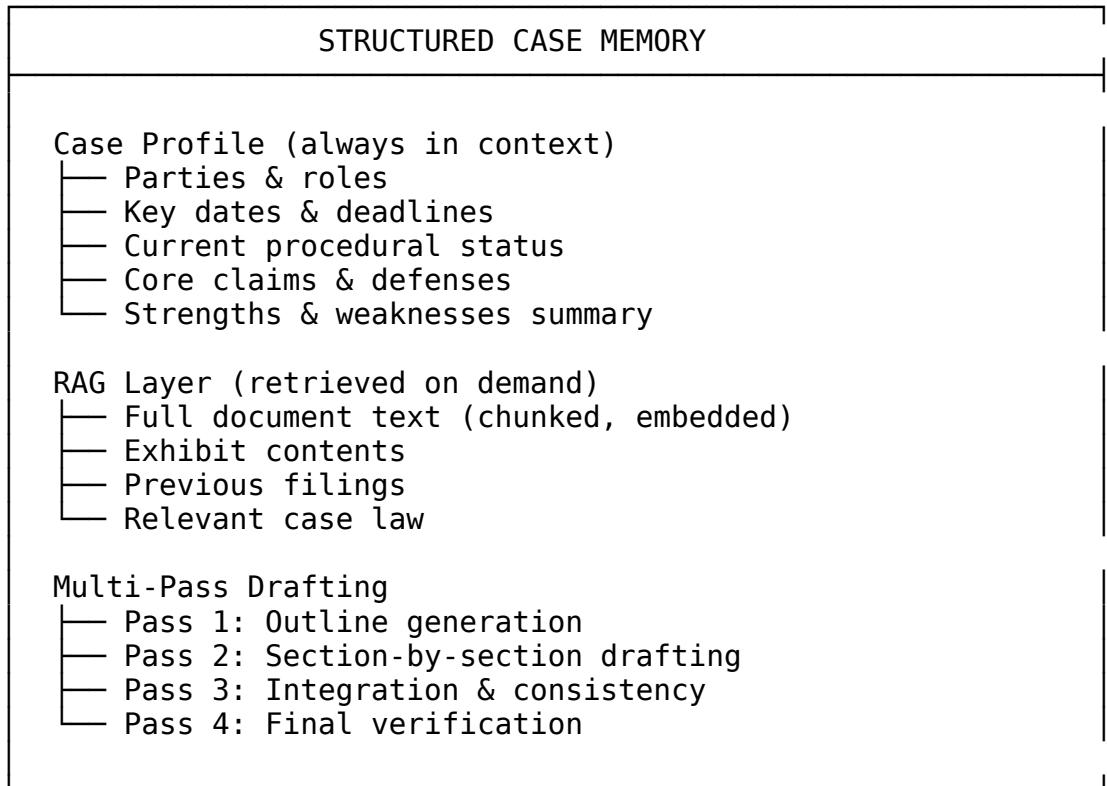
Context Window Strategy

Critical Insight: Large context ≠ legal reasoning quality

No model (open or closed) should ingest an entire lawsuit at once. The correct approach:

```
┌─────────────────────────────────────────────────────────┐ |
│            STRUCTURED CASE MEMORY                        │ |
├─────────────────────────────────────────────────────────┤ |
│                                                         │
│   Case Profile (always in context)                      │ |
│   ├── Parties & roles                                   │
│   ├── Key dates & deadlines                             │
│   ├── Current procedural status                         │
│   ├── Core claims & defenses                            │
│   └── Strengths & weaknesses summary                    │
│                                                         │ |
│   RAG Layer (retrieved on demand)                       │ |
│   ├── Full document text (chunked, embedded)            │
│   ├── Exhibit contents                                  │
│   ├── Previous filings                                  │
│   └── Relevant case law                                 │ |
│                                                         │
│   Multi-Pass Drafting                                   │ |
│   ├── Pass 1: Outline generation                        │
│   ├── Pass 2: Section-by-section drafting               │
│   ├── Pass 3: Integration & consistency                 │
│   └── Pass 4: Final verification                        │ |
│                                                         │ |
└─────────────────────────────────────────────────────────┘
```

Quality Tiers (User-Facing)

Tier	What User Sees	Backend Reality	Price
Standard	"Good for drafts"	Gemini + Sonnet	Free (limited)
Professional	"Recommended"	Gemini + Sonnet + verification	$0.10/doc
Premium	"Best quality"	Gemini + Opus + verification	$0.25/doc

Users never see model names. They see quality and use-case descriptions.

Provider Stack

Layer 1 & 3 (Cheap Operations): - Primary: Google Gemini 1.5 Flash - Fallback: OpenAI GPT-4o-mini - Cost: ~$0.075/1M input, $0.30/1M output

Layer 2 (Legal Reasoning): - Standard: Claude 3.5 Sonnet - Premium: Claude Opus 4 - Cost: Sonnet ~$3/1M input, $15/1M output - Cost: Opus ~$15/1M input, $75/1M output

Cost Model

Per-User Economics (Moderate Usage):

Component	Calls/Month	Cost/Call	Monthly
Ingestion (Gemini)	50	$0.002	$0.10
Reasoning (Sonnet)	30	$0.05	$1.50
Verification (Gemini)	30	$0.002	$0.06
Total COGS/user			**~$1.70**

Platform Economics:

Users	Monthly AI COGS	Infra	Total
100	$170	$50	$220
1,000	$1,700	$100	$1,800
10,000	$17,000	$500	$17,500

Compare to Self-Hosted: - Minimum viable GPU setup: $3,500/month - Breakeven vs API: ~2,000+ active users - And still inferior reasoning quality

Important Distinctions

Claude Max ≠ Claude API - Claude Max = subscription for claude.ai web usage - Claude API = separate billing, requires $5 minimum credit - They are completely separate products

Ollama/Local = Privacy Option Only - Not a cost savings - Not better quality - Useful for: users who cannot send data to cloud - We support it but don't recommend it

Model Routing Logic

```
function selectModel(task, tier, sensitivity) {
  // Layer 1: Always cheap
  if (task in ['ocr', 'extract', 'classify']) {
    return 'gemini-flash';
  }

  // Layer 2: Based on tier
  if (task in ['draft', 'analyze', 'strategize']) {
    if (tier === 'premium') return 'claude-opus';
    return 'claude-sonnet';
  }

  // Layer 3: Always cheap
  if (task in ['verify', 'check', 'validate']) {
    return 'gemini-flash';
```

```
  }

  // Privacy override
  if (sensitivity === 'high' && user.hasOllama) {
    return 'ollama-local';
  }
}
```

Future: When to Add Open Source

Open source becomes viable when: 1. **Volume**: 5,000+ active users 2. **Predictability**: Stable usage patterns 3. **Use case**: Bulk processing, not frontier reasoning 4. **Team**: Dedicated ML ops capability

Target open-source use cases: - Document OCR at scale - Embedding generation - Classification/routing - NOT legal reasoning or drafting

---

Jurisdictions

Phase 1: Hong Kong
- High Court Civil Procedure Rules
- District Court Rules
- Lands Tribunal
- Labour Tribunal

Phase 2: UK
- Civil Procedure Rules (CPR)
- Family Procedure Rules
- Employment Tribunals

Phase 3: US
- Federal Rules of Civil Procedure
- State-specific rules (California, New York, Texas priority)

Phase 4: Other Common Law
- Australia
- Canada
- Singapore

Privacy & Security

Data Handling Principles
1. **Client-side first** - All case data stored locally in browser/app
2. **No server storage** - We never store your documents or case details

3. **Encrypted in transit** - All API calls over HTTPS
4. **Minimal AI sharing** - Only send what's needed for each request
5. **Optional local AI** - Ollama option for complete privacy

What Gets Sent to AI Providers

Data Type	Sent to AI?	Purpose
Document text	Yes (when analyzing)	Analysis/generation
Case summary	Yes (for context)	Better responses
Personal details	Minimal	Only if in documents
API keys	Never	Stored locally only
Usage patterns	No	N/A

Compliance Considerations

- **GDPR**: No EU data storage, user controls all data
- **Attorney-Client Privilege**: Tool is not a lawyer, no privilege created
- **Court Rules**: Users responsible for accuracy of filings

Business Model Options

Option A: Freemium SaaS

```
FREE TIER
- 3 document generations/month
- Basic case profile
- Limited AI chat (10 messages/day)
- Community support

PRO ($29/month)
- Unlimited document generation
- Full case analysis
- Unlimited AI chat
- Email support
- Google integrations

TEAM ($99/month)
- Multiple cases
- Priority support
- API access
- Custom templates
```

Option B: One-Time Purchase

```
PERSONAL LICENSE ($99)
- Desktop app (Mac/Windows)
- Lifetime updates for 1 year
```

```
- All features unlocked
- Bring your own AI API key

PROFESSIONAL ($249)
- + 5 years updates
- + Priority support
- + Custom jurisdiction templates
```

Option C: Open Source + Services

```
OPEN SOURCE (FREE)
- Full app, self-hosted
- Bring your own AI key
- Community support

MANAGED SERVICE ($19/month)
- Hosted version
- Included AI credits
- Auto-updates
- Support

ENTERPRISE (Custom)
- Legal aid organizations
- Custom integrations
- Training
```

Competitive Landscape

Product	Strengths	Weaknesses	Our Advantage
LegalZoom	Brand recognition	Generic templates, expensive	Case-specific AI
Rocket Lawyer	Lawyer network	Subscription model, not AI-native	Deep case understanding
DoNotPay	Automation	Limited scope, US only	Comprehensive toolkit
ChatGPT	General capability	No case context, generic	Persistent case profile
Clio	Full practice management	For lawyers, not litigants	User-focused UX

Success Metrics

User Success

- Documents filed without rejection

- Cases won/settled favorably
- Time saved vs manual drafting
- User confidence scores

Product Metrics
- Daily active users
- Documents generated/user
- Feature adoption rates
- Session duration
- Retention (D1, D7, D30)

Business Metrics
- MRR/ARR
- Customer acquisition cost
- Lifetime value
- Churn rate
- NPS score

Roadmap

Q1 2025: Foundation
- ☒ Core document generation
- ☒ Case profile management
- ☒ Evidence management
- ☒ Timeline builder
- ☒ AI chat assistant
- ☐ Gemini API integration
- ☐ Polish onboarding flow

Q2 2025: Strategic Tools
- ☐ Enhanced case analyzer
- ☐ Motion response wizard
- ☐ Deadline automation
- ☐ Hearing simulator beta

Q3 2025: Integrations
- ☐ Gmail deep integration
- ☐ Calendar sync
- ☐ Google Drive backup
- ☐ PDF annotation

Q4 2025: Scale
- ☐ UK jurisdiction support

- • ☐ Mobile companion app
- • ☐ Legal aid partnerships
- • ☐ Community features

2026: Expansion
- • ☐ US state coverage
- • ☐ Video tutorials
- • ☐ Lawyer marketplace
- • ☐ API for integrations

Open Questions

1. **Pricing model** - Which model best serves access to justice mission while being sustainable?

2. **Jurisdiction priority** - Which markets after HK? UK has similar legal system, US has bigger market.

3. **AI provider default** - Gemini (free) vs Claude (quality)? Let user choose?

4. **Legal disclaimers** - How prominent? Every page or just onboarding?

5. **Community features** - Should users share templates? Privacy concerns?

6. **Lawyer involvement** - Review marketplace? Referral network?

Appendix: Document Types

Affidavit/Affirmation
- • Sworn statement of facts
- • Used to submit evidence
- • Requires proper jurat
- • Numbered paragraphs

Written Submissions
- • Legal arguments
- • Cites authorities
- • Responds to issues
- • Structured sections

Skeleton Argument
- • Brief argument outline
- • For hearings

- Key points only
- 3-5 pages max

Letter to Court
- Administrative requests
- Adjournments
- Directions
- Formal but concise

Response/Reply
- Answers to motions
- Point-by-point rebuttal
- Counter-arguments
- Prayer for relief

Document Version: 1.0 Last Updated: December 2024