

Represent Yourself - Product Specification

Executive Summary

Represent Yourself is an AI-powered legal assistant designed specifically for self-represented litigants (pro se / litigants in person). It combines document generation, case analysis, and procedural guidance into a native desktop application that helps ordinary people navigate complex legal proceedings without expensive legal representation.

Problem Statement

The Access to Justice Gap

1. **Legal representation is unaffordable** - Average hourly rates for lawyers range from \$200-\$500+. A typical civil case costs \$10,000-\$50,000+ in legal fees.
 2. **Self-representation is overwhelming** - Court procedures, legal terminology, document formatting, and deadlines create massive barriers for non-lawyers.
 3. **Existing tools are inadequate** - Legal templates are generic. ChatGPT gives generic answers. Neither understands your specific case context.
 4. **Information asymmetry** - Opposing parties with lawyers have massive advantages in knowing procedures, deadlines, and effective arguments.
-

Solution

Represent Yourself provides:

- **Deep case understanding** - Upload all your documents once, AI builds comprehensive case profile
- **Document generation** - Court-ready documents with proper formatting and legal language
- **Procedural guidance** - Step-by-step guidance through court procedures
- **Strategic analysis** - Honest assessment of case strengths, weaknesses, and risks
- **Deadline management** - Never miss a filing deadline

- **Hearing preparation** - Mock hearings and cross-examination practice
-

Target Users

Primary: Self-Represented Litigants

- **Civil litigation** defendants/plaintiffs
- **Family law** (divorce, custody, child support)
- **Employment disputes** (wrongful termination, discrimination)
- **Landlord-tenant** disputes
- **Small claims** escalated to higher courts
- **Consumer protection** cases

User Personas

1. Sarah - Divorce Litigant - Going through contested divorce - Spouse has lawyer, she doesn't - Needs help with financial disclosures, custody arrangements - Budget: Can't afford \$300/hour lawyer

2. Marcus - Employment Dispute - Wrongfully terminated, filed complaint - Company responded with motion to dismiss - Needs to file opposition within 14 days - Has evidence but doesn't know how to present it

3. Chen - Landlord Dispute - Suing landlord for security deposit - Case escalated beyond small claims - Intimidated by formal court procedures - English is second language

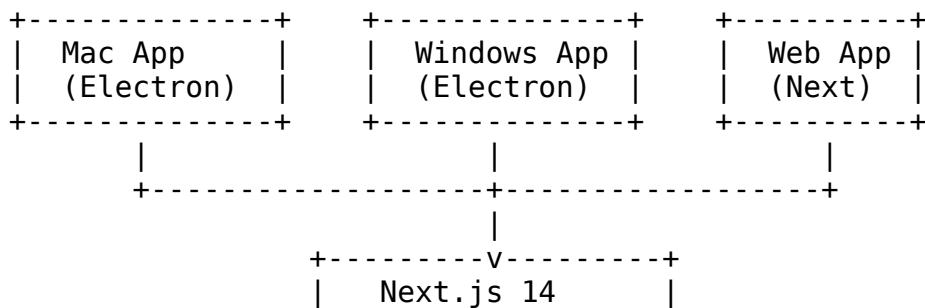
Secondary: Legal Aid Organizations

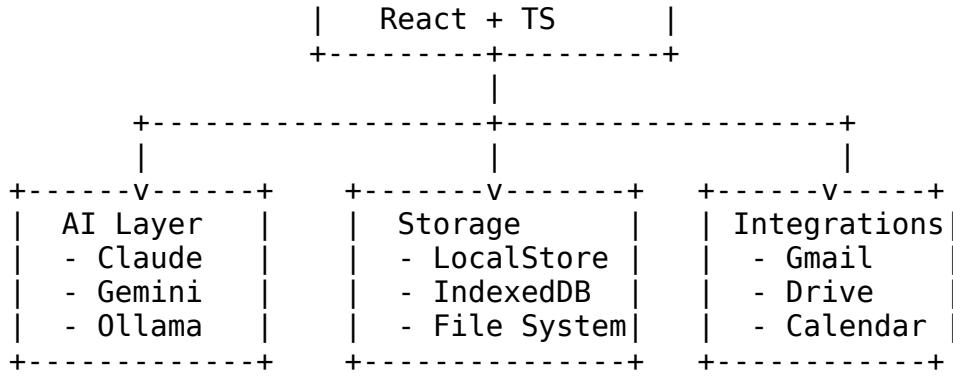
- Provide tool to clients they can't fully represent
 - Training resource for pro bono attorneys
 - Intake and case assessment assistance
-

Product Architecture

Platform Strategy

REPRESENT YOURSELF





Technical Stack

Layer	Technology	Purpose
Frontend	Next.js 14, React 18, TypeScript	App framework
Styling	Tailwind CSS, Radix UI	Design system
Desktop	Electron 28	Native Mac/Windows apps
State	Zustand + localStorage	Client-side persistence
AI	Claude API, Gemini API, Ollama	LLM providers
Integrations	Google APIs (Gmail, Drive, Calendar)	Productivity
PDF	pdf.js, pdf-lib	Document handling

Data Architecture

User Data (100% Client-Side)

```

+-- Case Profile
|   +-- Parties (plaintiffs, defendants, lawyers)
|   +-- Court details (case number, judge, jurisdiction)
|   +-- Procedural history
|   +-- Related proceedings

+-- Documents
|   +-- Uploaded files (PDF, DOCX, images)
|   +-- AI analysis results
|   +-- Generated drafts
|   +-- Filed documents

+-- Exhibits
|   +-- Exhibit metadata
|   +-- Bundles
|   +-- Page references

+-- Timeline

```

```
|   +-+ Events (filings, hearings, deadlines)
|   +-+ AI-extracted dates
|   +-+ User-added events
|
+-+ Configuration
  +-+ AI provider settings
  +-+ Google OAuth tokens
  +-+ UI preferences
```

Feature Set

Phase 1: Core (Current)

1. Case Setup & Onboarding

- AI-powered setup wizard
- Upload documents, AI extracts case details
- Guided questions to fill gaps
- Auto-generates case profile

2. Document Drafting

- **5 document types:**
 - Affidavit/Affirmation
 - Written Submissions
 - Skeleton Arguments
 - Letter to Court
 - Response/Reply
- Step-by-step wizard interface
- Court-specific formatting
- Export to DOCX/PDF

3. Document Analysis

- Upload any legal document
- AI extracts:
 - Document type and purpose
 - Key facts and dates
 - Legal issues raised
 - Required responses
 - Deadlines triggered

4. Evidence Management

- Exhibit cataloging with metadata
- Bundle creation for hearings
- Page numbering and indexing

- Status tracking (draft to filed to admitted)

5. Timeline Builder

- Chronological case visualization
- Event categorization (court/opposing/yours/victory)
- Deadline highlighting
- AI-assisted date extraction from documents

6. AI Assistant Chat

- Context-aware legal Q&A
- Knows your case details
- Explains procedures in plain English
- Suggests next steps

7. Plain Language Translator

- Paste legal jargon, get plain English
- Key term glossary
- “What this means for you” summary
- Action items extraction

Phase 2: Strategic Tools

8. Case Analyzer

- Comprehensive strength/weakness assessment
- Risk scoring (0-100)
- Win probability estimation
- Strategic recommendations

9. Procedure Guide

- Court-specific procedure handbook
- Step-by-step filing instructions
- Form requirements
- Fee schedules

10. Discovery Helper

- Analyze discovery requests
- Identify responsive documents
- Draft objections
- Privilege log assistance

11. Motion Response Wizard

- Parse opposing motions
- Generate point-by-point responses
- Legal authority suggestions

- Filing checklist

12. Deadline Calculator

- Auto-calculate procedural deadlines
- Calendar integration
- Reminder notifications
- “Days remaining” tracking

Phase 3: Advanced Features

13. Hearing Simulator

- Mock hearing practice
- AI plays opposing counsel
- Cross-examination prep
- Judge Q&A simulation

14. Opponent Intelligence

- Track opposing party patterns
- Analyze their filing history
- Predict likely arguments
- Identify weaknesses

15. Settlement Calculator

- Model settlement scenarios
- Risk-adjusted valuations
- Negotiation ranges
- Walk-away points

16. Risk Scorecard

- Multi-factor risk assessment
- Litigation cost projections
- Outcome probability matrix
- Decision support

17. Legal GPS

- Visual case navigation
- “You are here” indicator
- Next steps guidance
- Milestone tracking

18. Bundle Generator

- Auto-generate hearing bundles
- Court-compliant formatting
- Index generation

- Paginated PDF output

Phase 4: Integrations

19. Gmail Integration

- Monitor case-related emails
- Auto-categorize by sender
- Extract deadlines from emails
- Draft responses

20. Google Drive Integration

- Sync documents to Drive
- Collaborative editing
- Version history
- Backup protection

21. Calendar Integration

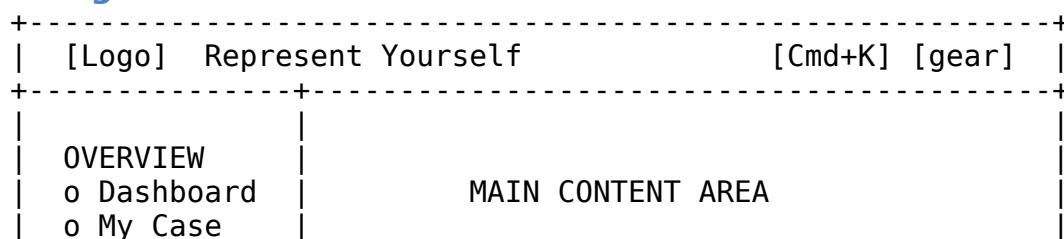
- Sync deadlines and hearings
 - Reminder scheduling
 - Preparation timelines
-

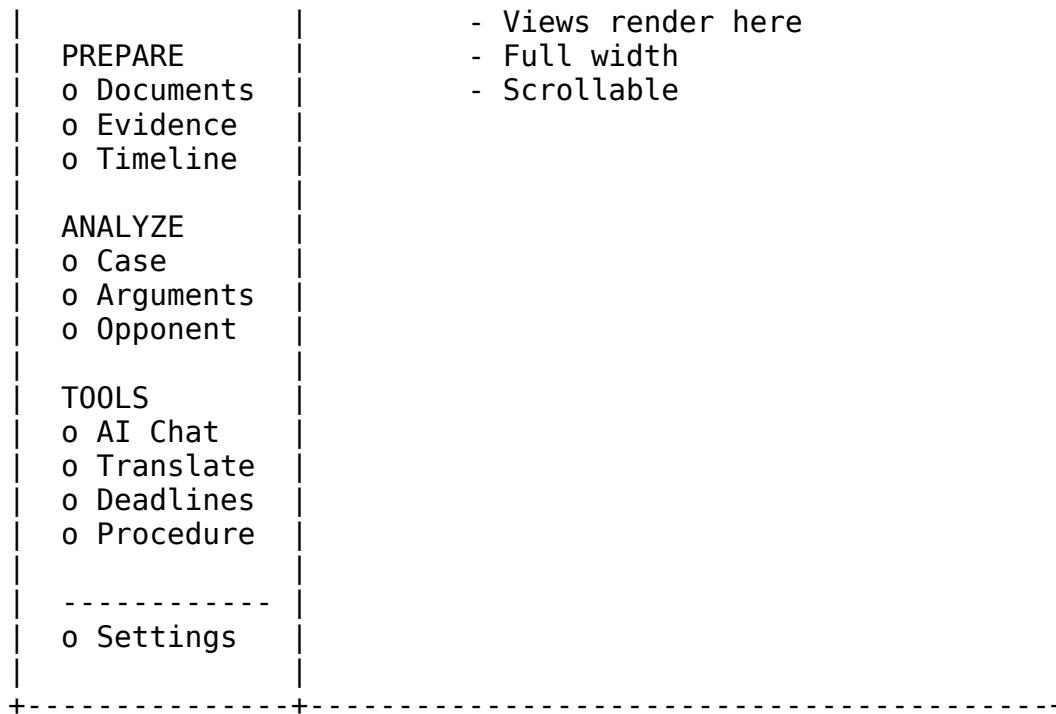
User Experience

Design Principles

1. **Calm, not overwhelming** - Legal matters are stressful. UI should feel supportive, not add anxiety.
2. **Guide, don't dump** - Never show 50 options at once. Progressive disclosure with clear next steps.
3. **Plain language first** - Explain everything in plain English. Legal terms always have definitions.
4. **Mobile-ready but desktop-focused** - Legal work needs big screens. Mobile for reference only.
5. **Offline-capable** - Documents and case data work offline. AI features need connection.

Navigation Structure

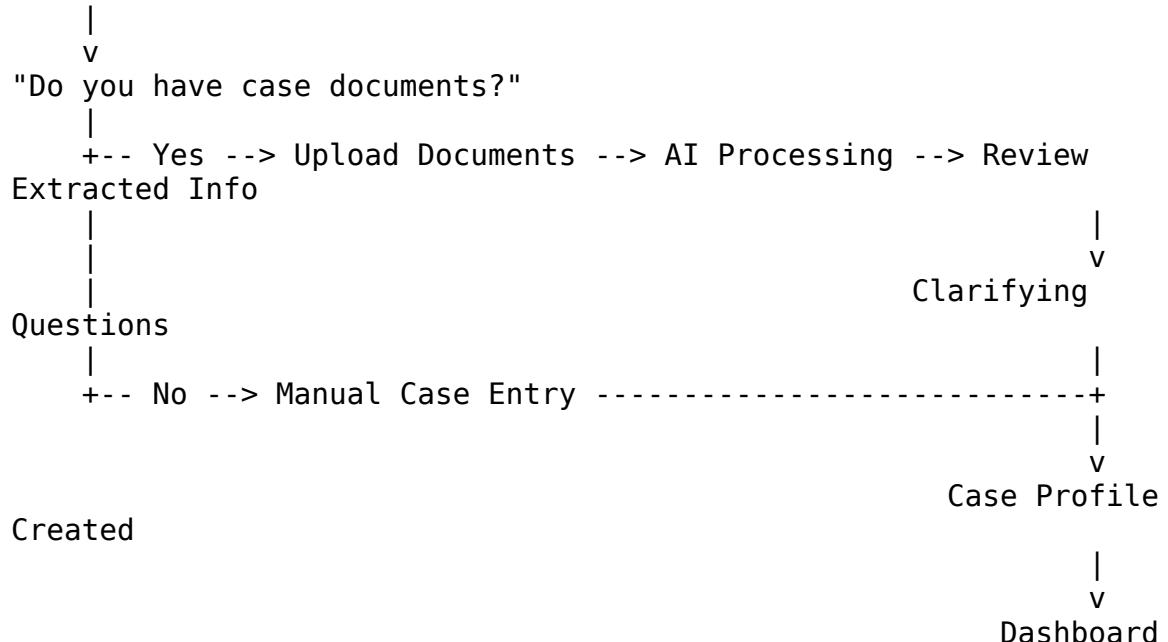




Key Flows

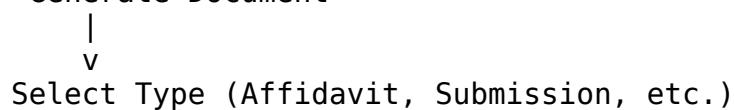
Flow 1: First-Time Setup

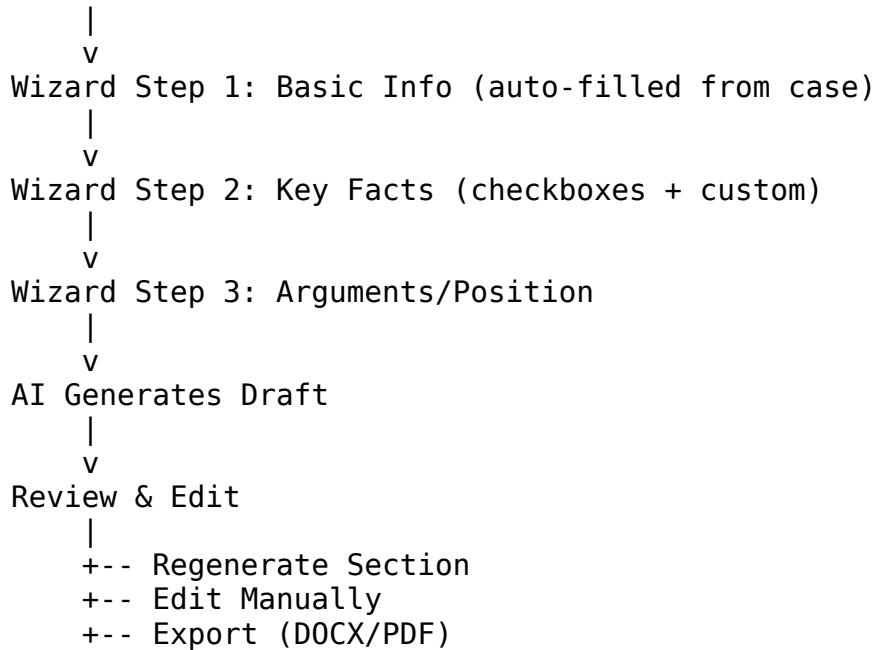
Welcome Screen



Flow 2: Document Generation

"Generate Document"





AI Strategy

Strategic Foundation

Core Decision: API-First, Not Open Source

Open-source LLMs are a margin-expansion tool, not a foundation:

- At low user counts, APIs are significantly cheaper
- Fixed infrastructure costs for self-hosted: \$3,500-8,000+/month
- API costs at early stage: \$400-900/month
- Open-source only becomes economical at large, predictable scale

Why Not Self-Hosted Open Source:

- High fixed infrastructure costs
- DevOps + ML maintenance burden
- Lower reasoning reliability for legal tasks
- Silent failures are unacceptable in legal products
- No frontier open model can run locally (Qwen3-235B needs 4-8x H100 GPUs)

Three-Layer Model Architecture

Users choose **quality tier**, not models. The system routes intelligently.

USER QUALITY SELECTION		
[Standard]	[Professional]	[Premium]
Free	\$0.10/doc	\$0.25/doc



LAYER 1: INGESTION & STRUCTURING
(Cheap Model)

- OCR and text extraction
- Timeline extraction
- Entity recognition (parties, dates, amounts)
- Document classification
- Metadata extraction

Model: Gemini Flash / GPT-4o-mini

Cost: ~\$0.001-0.005 per document



LAYER 2: LEGAL REASONING & DRAFTING
(Premium Model)

- Document generation (affidavits, submissions)
- Case strength analysis
- Argument identification
- Strategic recommendations
- Motion response drafting

Standard: Claude Sonnet (default)

Premium: Claude Opus (high-stakes toggle)

Cost: \$0.01-0.10 per generation



LAYER 3: VERIFICATION & CONSISTENCY
(Cheap Model)

- Citation checking
- Internal consistency validation
- Format compliance
- Deadline verification
- Cross-reference validation

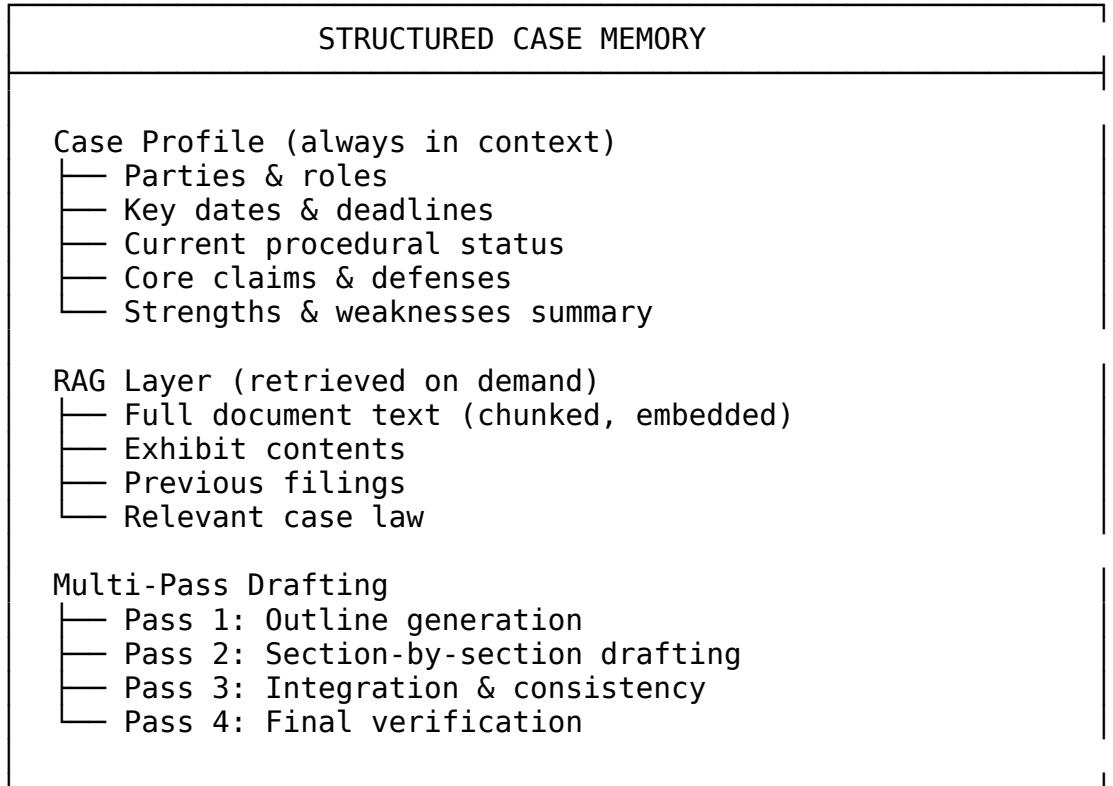
Model: Gemini Flash / GPT-4o-mini

Cost: ~\$0.001-0.005 per check

Context Window Strategy

Critical Insight: Large context ≠ legal reasoning quality

No model (open or closed) should ingest an entire lawsuit at once. The correct approach:



Quality Tiers (User-Facing)

Tier	What User Sees	Backend Reality	Price
Standard	“Good for drafts”	Gemini + Sonnet	Free (limited)
Professional	“Recommended”	Gemini + Sonnet + verification	\$0.10/doc
Premium	“Best quality”	Gemini + Opus + verification	\$0.25/doc

Users never see model names. They see quality and use-case descriptions.

Provider Stack

Layer 1 & 3 (Cheap Operations): - Primary: Google Gemini 1.5 Flash - Fallback: OpenAI GPT-4o-mini - Cost: ~\$0.075/1M input, \$0.30/1M output

Layer 2 (Legal Reasoning): - Standard: Claude 3.5 Sonnet - Premium: Claude Opus 4 - Cost: Sonnet ~\$3/1M input, \$15/1M output - Cost: Opus ~\$15/1M input, \$75/1M output

Cost Model

Per-User Economics (Moderate Usage):

Component	Calls/Month	Cost/Call	Monthly
Ingestion (Gemini)	50	\$0.002	\$0.10
Reasoning (Sonnet)	30	\$0.05	\$1.50
Verification (Gemini)	30	\$0.002	\$0.06
Total COGS/user	~\$1.70		

Platform Economics:

Users	Monthly AI COGS	Infra	Total
100	\$170	\$50	\$220
1,000	\$1,700	\$100	\$1,800
10,000	\$17,000	\$500	\$17,500

Compare to Self-Hosted: - Minimum viable GPU setup: \$3,500/month - Breakeven vs API: ~2,000+ active users - And still inferior reasoning quality

Important Distinctions

Claude Max ≠ Claude API - Claude Max = subscription for claude.ai web usage - Claude API = separate billing, requires \$5 minimum credit - They are completely separate products

Ollama/Local = Privacy Option Only - Not a cost savings - Not better quality - Useful for: users who cannot send data to cloud - We support it but don't recommend it

Model Routing Logic

```
function selectModel(task, tier, sensitivity) {
    // Layer 1: Always cheap
    if (task in ['ocr', 'extract', 'classify']) {
        return 'gemini-flash';
    }

    // Layer 2: Based on tier
    if (task in ['draft', 'analyze', 'strategize']) {
        if (tier === 'premium') return 'claude-opus';
        return 'claude-sonnet';
    }

    // Layer 3: Always cheap
    if (task in ['verify', 'check', 'validate']) {
        return 'gemini-flash';
    }
}
```

```
}

// Privacy override
if (sensitivity === 'high' && user.hasOllama) {
  return 'ollama-local';
}
}
```

Future: When to Add Open Source

Open source becomes viable when: 1. **Volume**: 5,000+ active users 2. **Predictability**: Stable usage patterns 3. **Use case**: Bulk processing, not frontier reasoning 4. **Team**: Dedicated ML ops capability

Target open-source use cases: - Document OCR at scale - Embedding generation - Classification/routing - NOT legal reasoning or drafting

Jurisdictions

Phase 1: Hong Kong

- High Court Civil Procedure Rules
- District Court Rules
- Lands Tribunal
- Labour Tribunal

Phase 2: UK

- Civil Procedure Rules (CPR)
- Family Procedure Rules
- Employment Tribunals

Phase 3: US

- Federal Rules of Civil Procedure
- State-specific rules (California, New York, Texas priority)

Phase 4: Other Common Law

- Australia
 - Canada
 - Singapore
-

Privacy & Security

Deployment Options

Two ways to use Represent Yourself - user chooses based on trust level:

DEPLOYMENT OPTIONS

OPTION A: CLOUD (app.representyourself.legal)

- Hosted by us, managed by us
- Zero setup - just sign up
- We handle AI API keys (included in price)
- Data passes through our servers
- For users who value convenience

Trust model: "I trust you with my data"

Price: Subscription tiers (\$50-\$750/mo)

OPTION B: SELF-HOSTED (download + run locally)

- Runs on user's machine (Electron app or Docker)
- User provides own API keys (BYOK)
- Data never leaves their device
- API calls go directly to providers
- For privacy-conscious users

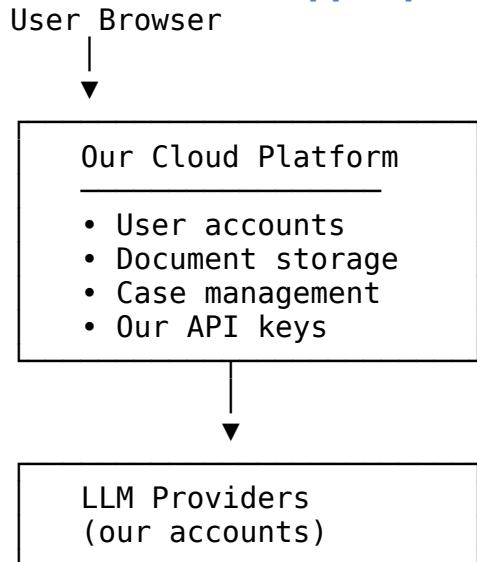
Trust model: "I trust Anthropic/Google, not you"

Price: License fee (\$25/mo) + user pays own API costs

Architecture Comparison

Aspect	Cloud	Self-Hosted
Hosting	Our servers	User's machine
Setup	Sign up, done	Download + configure
AI Keys	We provide	User provides (BYOK)
Data Flow	Through our servers	Direct to providers
Updates	Automatic	Manual or auto-update
Support	Full support	Community + docs
Price	\$50-750/mo (all-in)	\$25/mo + API costs
Privacy	Standard	Maximum

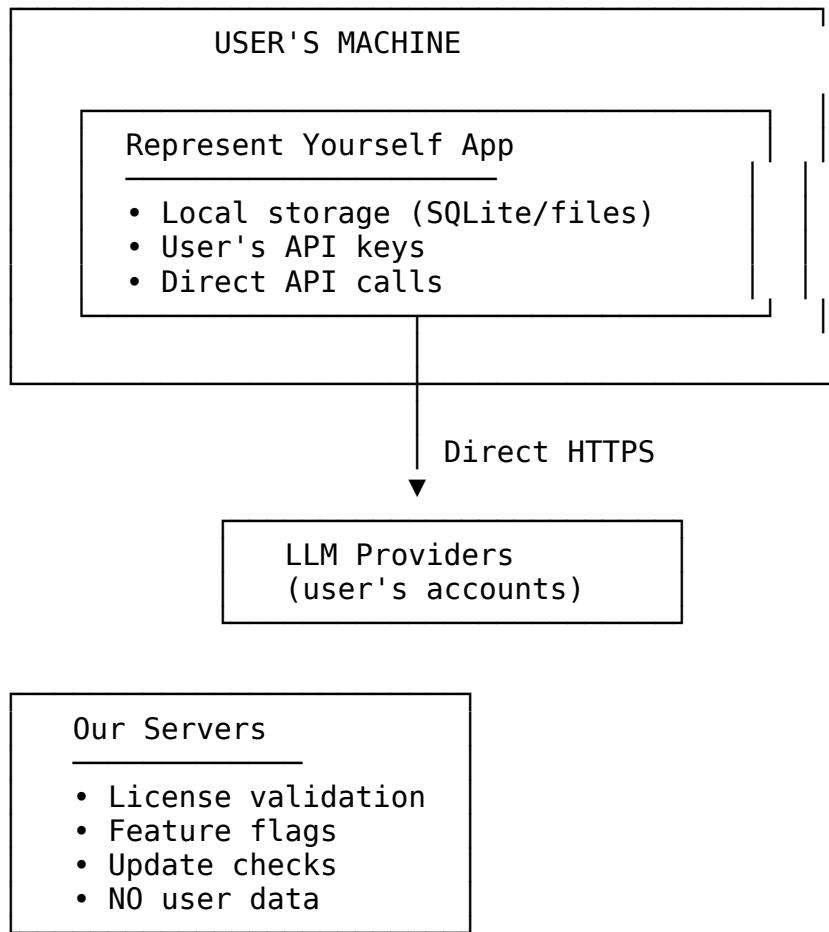
Cloud Version (app.representyourself.legal)



Pros: - Zero friction signup - No technical setup - Included AI credits - Cross-device sync - Automatic backups

Cons: - Data on our servers - Dependent on our uptime - Less privacy

Self-Hosted Version (Desktop App / Docker)



Pros: - Maximum privacy - Data never leaves device - User controls everything - Works offline (except AI) - No vendor lock-in

Cons: - Technical setup required - User manages own API keys - No cross-device sync - Manual backups

Self-Hosted Deployment Options

Option B1: Desktop App (Electron)

```
# Download from releases page  
# Double-click to install  
# Enter license key  
# Add your API keys in settings
```

- Mac (Intel + Apple Silicon)
- Windows
- Linux

Option B2: Docker (for servers/NAS)

```
docker run -d \
-p 3000:3000 \
-v ~/represent-data:/data \
-e LICENSE_KEY=xxx \
representyourself/app:latest
```

- Run on home server
- Run on NAS (Synology, etc.)
- Run on VPS (if user wants own cloud)

Option B3: Source Code (developers)

```
git clone https://github.com/representyourself/app
npm install
npm run build
npm start
```

- Full customization
- Self-audit the code
- Contribute improvements

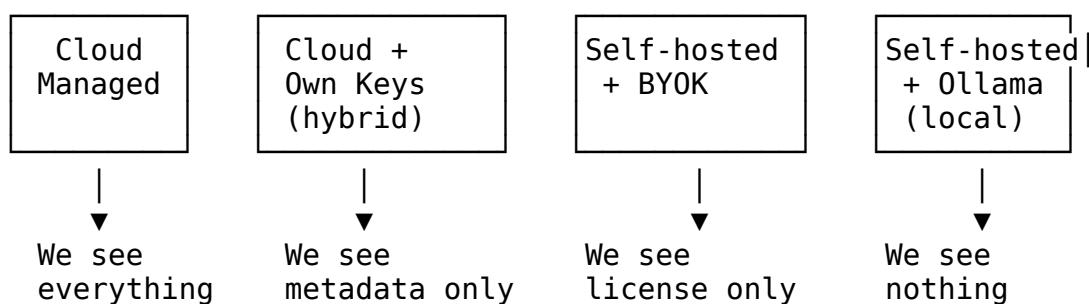
Pricing by Deployment

Plan	Cloud	Self-Hosted
Starter	\$50/mo (2 docs included)	\$25/mo + own API
Pro	\$150/mo (6 docs included)	\$25/mo + own API
Litigation	\$350/mo (15 docs included)	\$25/mo + own API
High-Stakes	\$750/mo (premium AI included)	\$50/mo + own API

Self-hosted users pay flat license fee + their own API costs. Typical API cost: \$5-30/month depending on usage.

Trust Spectrum

MAXIMUM CONVENIENCE ← → MAXIMUM PRIVACY



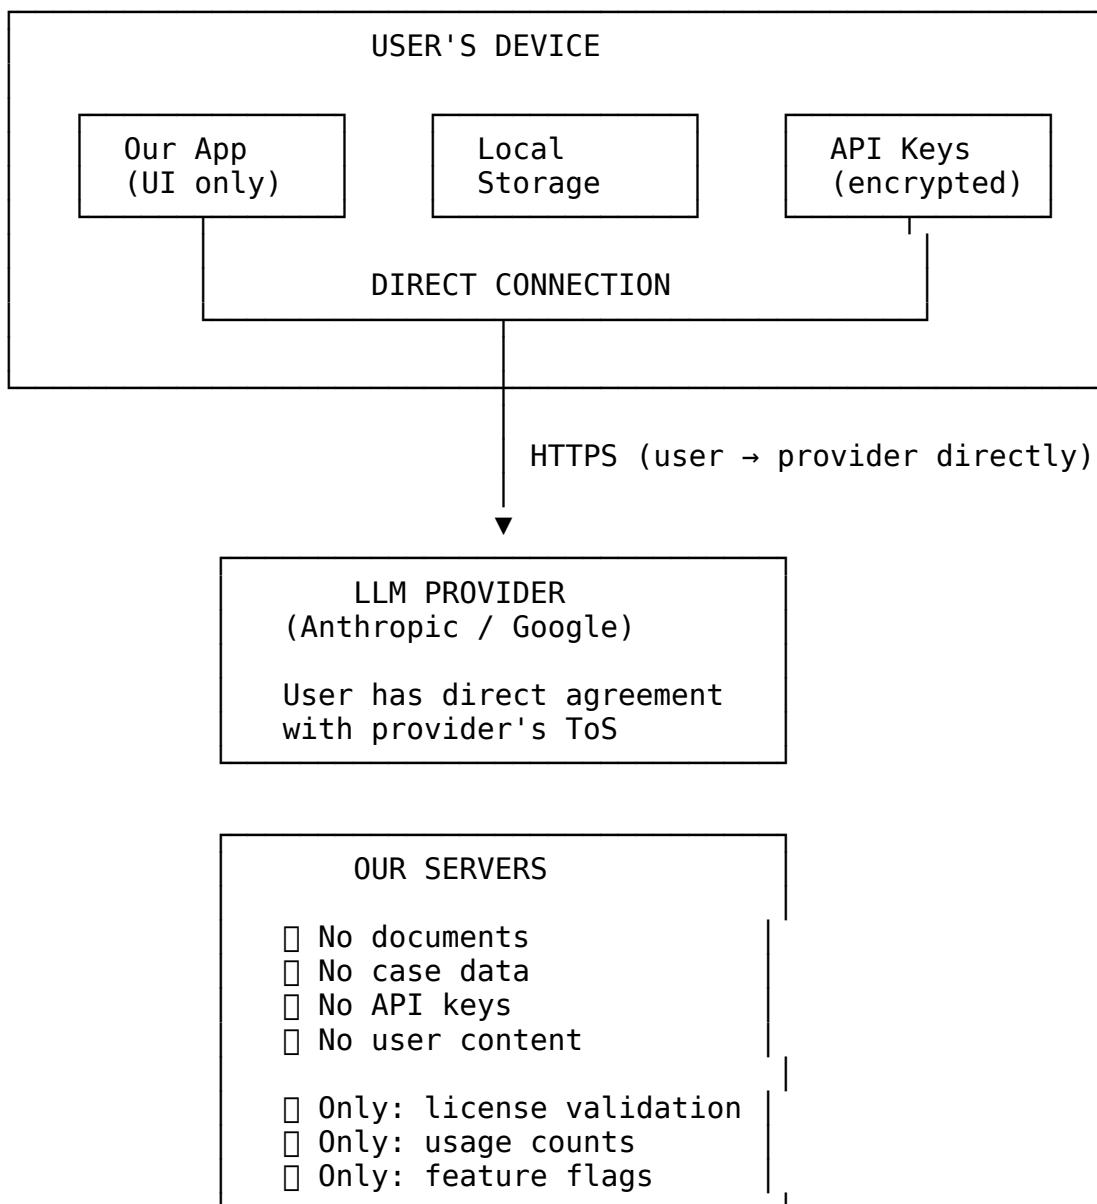
The Trust Problem (Addressed)

User A: "I just want it to work." → Use Cloud version. We handle everything.

User B: "I don't want you seeing my documents." → Use Self-hosted + BYOK. We only validate your license.

User C: "I don't trust anyone with my data." → Use Self-hosted + Ollama. Nothing leaves your machine.

Zero-Knowledge Architecture



What We See vs What Providers See

Data	Us (Represent Yourself)	LLM Provider
Documents	<input type="checkbox"/> Never	<input type="checkbox"/> When analyzing
Case details	<input type="checkbox"/> Never	<input type="checkbox"/> When generating
Personal info	<input type="checkbox"/> Never	<input type="checkbox"/> If in documents
API keys	<input type="checkbox"/> Never	<input type="checkbox"/> User's own key
Usage count	<input type="checkbox"/> Aggregate only	<input type="checkbox"/> Per-request
Feature usage	<input type="checkbox"/> Anonymous	<input type="checkbox"/> No

Trust Model Options

Option 1: BYOK (Maximum Privacy)

User provides own API keys

- Calls go directly from their device to provider
- We see: nothing about their content
- Provider sees: their requests (under user's account)
- User's existing ToS with provider applies

Option 2: Managed Keys (Convenience)

We provide API access (pooled keys)

- Calls still go from their device
- We see: usage counts, not content
- Provider sees: requests under our account
- Our ToS + provider ToS apply

Option 3: Local Only (Paranoid Mode)

Ollama on user's machine

- Nothing leaves their device
- We see: nothing
- Provider sees: nothing
- User responsible for model quality

Implementation: Client-Side Direct Calls

All LLM API calls are made directly from the user's browser/app:

```
// This happens ON THE USER'S DEVICE
async function callLLM(prompt: string) {
    // API key from user's local storage
    const apiKey = localStorage.getItem('user_api_key');

    // Direct call to provider - never touches our servers
    const response = await
fetch('https://api.anthropic.com/v1/messages', {
    headers: {
```

```

    'x-api-key': apiKey, // User's own key
    'anthropic-dangerous-direct-browser-access': 'true'
  },
  body: JSON.stringify({ messages: [...] })
);

// Response stays on user's device
return response.json();
}

```

Key point: The anthropic-dangerous-direct-browser-access header enables browser-to-API calls without a backend proxy. Same pattern works for Gemini.

What Our Backend Does (Minimal)

Our API handles ONLY:

- License key validation (is subscription active?)
- Document count tracking (how many this month?)
- Feature flag checks (which tier features enabled?)
- Anonymous analytics (which features are popular?)

Our API NEVER receives:

- Document content
- Case details
- User's API keys
- Generated outputs
- Chat histories

Data Residency Guarantees

Data Type	Location	Encrypted	We Can Access?
Documents	User's device only	Yes (browser)	<input type="checkbox"/> No
Case profile	User's device only	Yes (browser)	<input type="checkbox"/> No
API keys	User's device only	Yes (encrypted)	<input type="checkbox"/> No
Chat history	User's device only	Yes (browser)	<input type="checkbox"/> No
Usage counts	Our servers	Yes	<input type="checkbox"/> Aggregate
Subscription	Our servers	Yes	<input type="checkbox"/> Yes

Privacy Comparison

Approach	Privacy	Convenience	Cost
BYOK	Maximum	Medium	User pays provider

Approach	Privacy	Convenience	Cost
Managed	High	Maximum	We pay provider
Local Ollama	Absolute	Low	Free (GPU costs)

User-Facing Privacy Controls

Settings → Privacy

- Use my own API keys (BYOK mode)
- Enable local processing only (Ollama)
- Clear all local data
- Export my data (JSON)
- Delete my account

Compliance Considerations

- **GDPR:** User controls all PII, we process nothing
- **HIPAA:** Not applicable (we don't store PHI)
- **Attorney-Client Privilege:** Tool is not a lawyer, no privilege created
- **Court Rules:** Users responsible for accuracy of filings
- **Data Sovereignty:** All data stays on user's device

The Privacy Pitch

“Your legal documents never leave your device. We’re just the UI layer. When you use AI features, your device talks directly to Anthropic or Google - under your own account, your own API keys, your own agreement with them. We couldn’t read your documents even if we wanted to.”

Business Model

Pricing Strategy: Document-Based Subscriptions

Pricing based on document generation, not seats or features. Clear value exchange.

Subscription Tiers (USD)

Starter — \$50/month

For: Self-represented litigants, light matters, first-time users

Included	Overage
2 standard documents/month	\$35/document
Basic case profile	
AI chat (limited)	

Included	Overage
Email support	

Economics:

Revenue:	\$50
COGS ($2 \times \$9$):	\$18
Gross Margin:	64%

This tier converts skeptics and filters unserious users.

Pro — \$150/month

For: Active litigants, ongoing cases, regular filings

Included	Overage
6 standard documents/month	\$30/document
Full case profile	
Unlimited AI chat	
Priority queue	
Evidence management	

Economics:

Revenue:	\$150
COGS ($6 \times \$9$):	\$54
Gross Margin:	64%

Core revenue tier. Most users land here.

Litigation — \$350/month

For: Heavy users, multiple motions, complex cases

Included	Overage
15 standard documents/month	\$25/document
Extended context windows	
Deep document retrieval	
Long-form pleadings enabled	
Bundle generation	
Timeline automation	

Economics:

Revenue: \$350
 COGS ($15 \times \$9$): \$135
 Gross Margin: 61%

High volume, still healthy margins.

High-Stakes — \$750/month

For: Serious litigation, committal, summary judgment, appeals

Included	Overage
10 standard documents/month	\$30/document
5 high-stakes documents (Opus)	\$120/document
Maximum reasoning depth	
Multi-pass verification	
Hearing simulator	
Opponent intelligence	
Direct support channel	

Economics:

Revenue: \$750
 COGS:
 10 standard $\times \$9$: \$90
 5 high-stakes $\times \$18$: \$90
 Total COGS: \$180
 Gross Margin: 76%

This tier prints money and signals seriousness.

Tier Comparison

Feature	Starter	Pro	Litigation	High-Stakes
Price	\$50	\$150	\$350	\$750
Standard Docs	2	6	15	10
High-Stakes Docs	-	-	-	5
Overage (std)	\$35	\$30	\$25	\$30

Feature	Starter	Pro	Litigation	High-Stakes
Overage (premium)	-	-	-	\$120
AI Chat	Limited	Unlimited	Unlimited	Unlimited
Context Depth	Basic	Standard	Extended	Maximum
Verification	-	-	Basic	Full
Support	Email	Email	Priority	Direct
Gross Margin	64%	64%	61%	76%

Unit Economics

Per-Document COGS Breakdown:

Component	Standard Doc	High-Stakes Doc
Ingestion (Gemini)	\$0.50	\$0.50
Reasoning (Sonnet)	\$6.00	-
Reasoning (Opus)	-	\$15.00
Verification	\$0.50	\$1.50
Infrastructure	\$2.00	\$1.00
Total COGS	\$9.00	\$18.00

Revenue Projections

Scenario	Users	Mix	MRR	Annual
Seed	50	60% Starter, 30% Pro, 10% Lit	\$4,250	\$51K
Early	200	40% Starter, 40% Pro, 15% Lit, 5% High	\$22,750	\$273K
Growth	1,000	30% Starter, 40% Pro, 20% Lit, 10%	\$152,500	\$1.83M

Scenario	Users	Mix	MRR	Annual
		High		

Alternative: BYOK (Bring Your Own Key)

For price-sensitive users or those with existing API credits:

BYOK Plan — \$25/month - All features unlocked - User provides own API keys (Claude, Gemini) - We charge for platform only - No included documents - ~90% margin on subscription

Good for developers, power users, privacy-conscious.

Competitive Landscape

Product	Strengths	Weaknesses	Our Advantage
LegalZoom	Brand recognition	Generic templates, expensive	Case-specific AI
Rocket Lawyer	Lawyer network	Subscription model, not AI-native	Deep case understanding
DoNotPay	Automation	Limited scope, US only	Comprehensive toolkit
ChatGPT	General capability	No case context, generic	Persistent case profile
Clio	Full practice management	For lawyers, not litigants	User-focused UX

Success Metrics

User Success

- Documents filed without rejection
- Cases won/settled favorably
- Time saved vs manual drafting
- User confidence scores

Product Metrics

- Daily active users
- Documents generated/user
- Feature adoption rates
- Session duration
- Retention (D1, D7, D30)

Business Metrics

- MRR/ARR
 - Customer acquisition cost
 - Lifetime value
 - Churn rate
 - NPS score
-

Roadmap

Q1 2025: Foundation

- Core document generation
- Case profile management
- Evidence management
- Timeline builder
- AI chat assistant
- Gemini API integration
- Polish onboarding flow

Q2 2025: Strategic Tools

- Enhanced case analyzer
- Motion response wizard
- Deadline automation
- Hearing simulator beta

Q3 2025: Integrations

- Gmail deep integration
- Calendar sync
- Google Drive backup
- PDF annotation

Q4 2025: Scale

- UK jurisdiction support
- Mobile companion app
- Legal aid partnerships
- Community features

2026: Expansion

- US state coverage
 - Video tutorials
 - Lawyer marketplace
 - API for integrations
-

Open Questions

1. **Pricing model** - Which model best serves access to justice mission while being sustainable?
 2. **Jurisdiction priority** - Which markets after HK? UK has similar legal system, US has bigger market.
 3. **AI provider default** - Gemini (free) vs Claude (quality)? Let user choose?
 4. **Legal disclaimers** - How prominent? Every page or just onboarding?
 5. **Community features** - Should users share templates? Privacy concerns?
 6. **Lawyer involvement** - Review marketplace? Referral network?
-

Appendix: Document Types

Affidavit/Affirmation

- Sworn statement of facts
- Used to submit evidence
- Requires proper jurat
- Numbered paragraphs

Written Submissions

- Legal arguments
- Cites authorities
- Responds to issues
- Structured sections

Skeleton Argument

- Brief argument outline
- For hearings
- Key points only
- 3-5 pages max

Letter to Court

- Administrative requests
- Adjournments
- Directions
- Formal but concise

Response/Reply

- Answers to motions

- Point-by-point rebuttal
 - Counter-arguments
 - Prayer for relief
-

Document Version: 1.0 Last Updated: December 2024